

The Dummy Variable Trap with R

Econ 440 - Introduction to Econometrics

Patrick Toche, ptoche@fullerton.edu

12 May 2022

Create regions:

To illustrate regression with binary variables, let's group California counties by geographical area, based on Census data. To keep the example simple, let's **arbitrarily** create 4 regions named: North, South, East, West, based on a cursory look at a map.

For instance, here is what the list *North.Counties* looks like.

North.Counties

```
## [1] "Butte"          "Colusa"          "El Dorado"       "Glenn"
## [5] "Lassen"         "Modoc"           "Nevada"          "Placer"
## [9] "Plumas"         "Sacramento"      "Shasta"          "Sierra"
## [13] "Siskiyou"       "Sutter"          "Tehama"          "Yolo"
## [17] "Yuba"           "Del Norte"       "Humboldt"        "Lake"
## [21] "Mendocino"      "Napa"            "Sonoma"          "Trinity"
## [25] "Alameda"        "Contra Costa"    "Marin"           "San Francisco"
## [29] "San Mateo"      "Santa Clara"     "Solano"          "Alpine"
## [33] "Amador"         "Calaveras"       "Madera"          "Mariposa"
## [37] "Merced"         "Mono"            "San Joaquin"     "Stanislaus"
## [41] "Tuolumne"
```

To view the complete code, see the Rmd file.

Create dummy variables for each region:

```
df$North <- df$county %in% North.Counties
df$East <- df$county %in% East.Counties
df$South <- df$county %in% South.Counties
df$West <- df$county %in% West.Counties
head(df[c("North", "East", "South", "West")], 10)
```

```
## # A tibble: 10 x 4
##   North East  South West
##   <lgl> <lgl> <lgl> <lgl>
## 1 TRUE  FALSE FALSE FALSE
## 2 TRUE  FALSE FALSE FALSE
## 3 TRUE  FALSE FALSE FALSE
## 4 TRUE  FALSE FALSE FALSE
## 5 TRUE  FALSE FALSE FALSE
## 6 FALSE TRUE  FALSE FALSE
## 7 TRUE  FALSE FALSE FALSE
## 8 FALSE TRUE  FALSE FALSE
```

```
## 9 FALSE TRUE FALSE FALSE
## 10 TRUE FALSE FALSE FALSE
```

Check consistency of the categories:

```
unique(df$North + df$East + df$South + df$West) == 1
```

```
## [1] TRUE
```

Create a categorical variable for all 4 regions:

```
df$Region <- NA
dummies <- c("North", "East", "South", "West")
for (col in dummies)
  df$Region[which(df[,col] == TRUE)] <- col
head(df$Region, 10)
```

```
## [1] "North" "North" "North" "North" "North" "East" "North" "East" "East"
## [10] "North"
```

Regression With Categorical Variables

If you estimate a linear regression with all categories and an intercept, R will automatically drop one of the categories:

```
m1 <- lm(TestScore ~ STR + Region, data=df)
summary(m1)
```

```
##
## Call:
## lm(formula = TestScore ~ STR + Region, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -52.03 -13.10  -0.82  12.64  45.42
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   671.079     9.965   67.34 < 2e-16 ***
## STR           -1.558     0.488   -3.20  0.0015 **
## RegionNorth    18.081     2.393    7.56  2.7e-13 ***
## RegionSouth    11.892     2.861    4.16  3.9e-05 ***
## RegionWest     17.054     3.461    4.93  1.2e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 17.4 on 415 degrees of freedom
## Multiple R-squared:  0.17, Adjusted R-squared:  0.162
## F-statistic: 21.2 on 4 and 415 DF, p-value: 6.38e-16
```

The categorical variable *Region* stores the categories as strings (aka characters). When running the regression, R temporarily creates factor variables using alphabetical ordering, thus creating categories for *RegionNorth*, *RegionSouth* and *RegionWest*, while omitting the first category in the list, *RegionEast*.

To use, say, *North*, as the reference region, create factors in the desired order, and run the regression again:

```
df$Region2 <- factor(df$Region, levels=c("North", "East", "South", "West"))
m2 <- lm(TestScore ~ STR + Region2, data=df)
summary(m2)
```

```
##
## Call:
## lm(formula = TestScore ~ STR + Region2, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -52.03 -13.10  -0.82   12.64   45.42
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   689.160     9.349   73.72 < 2e-16 ***
## STR           -1.558     0.488   -3.20  0.0015 **
## Region2East  -18.081     2.393   -7.56 2.7e-13 ***
## Region2South  -6.189     2.442   -2.53  0.0116 *
## Region2West   -1.027     3.072   -0.33  0.7383
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 17.4 on 415 degrees of freedom
## Multiple R-squared:  0.17,    Adjusted R-squared:  0.162
## F-statistic: 21.2 on 4 and 415 DF,  p-value: 6.38e-16
```

If the categories are already stored as factors but do not appear in the desired order, you can reset the factor levels:

```
df$Region3 <- relevel(df$Region2, ref="South")
m3 <- lm(TestScore ~ STR + Region3, data=df)
summary(m3)
```

```
##
## Call:
## lm(formula = TestScore ~ STR + Region3, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -52.03 -13.10  -0.82   12.64   45.42
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   682.971     10.349   65.99 < 2e-16 ***
## STR           -1.558     0.488   -3.20  0.0015 **
## Region3North    6.189     2.442    2.53  0.0116 *
## Region3East   -11.892     2.861   -4.16 3.9e-05 ***
## Region3West     5.162     3.430    1.50  0.1331
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 17.4 on 415 degrees of freedom
## Multiple R-squared:  0.17,    Adjusted R-squared:  0.162
## F-statistic: 21.2 on 4 and 415 DF,  p-value: 6.38e-16
```